Philosophical Studies
https://doi.org/10.1007/s11098-025-02379-9

# Counter-productivity and suspicion: two arguments against talking about the AGI control problem

**Jakob Stenseke**[1]

Accepted: 29 June 2025
© The Author(s) 2025

**Abstract**
How do you control a superintelligent artificial being given the possibility that its goals or actions might conflict with human interests? Over the past few decades, this concern– *the AGI control problem*– has remained a central challenge for research in AI safety. This paper develops and defends two arguments that provide pro tanto support for the following policy for those who worry about the AGI control problem: *don't talk about it.* The first is *argument from counter-productivity*, which states that unless kept secret, efforts to solve the control problem could be used by a misaligned AGI to counter those very efforts. The second is *argument from suspicion*, stating that open discussions of the control problem may serve to make humanity appear threatening to an AGI, which increases the risk that the AGI perceives humanity as a threat. I consider objections to the arguments and find them unsuccessful. Yet, I also consider objections to the don't-talk policy itself and find it inconclusive whether it should be adopted. Additionally, the paper examines whether the arguments extend to other areas of AI safety research, such as *AGI alignment*, and argues that they likely do, albeit not necessarily as directly. I conclude by offering recommendations on what one can safely talk about, regardless of whether the don't-talk policy is ultimately adopted.

**Keywords** AI alignment · AI control · Existential risk · Superintelligence · Artificial general intelligence

✉ Jakob Stenseke
Jakob.stenseke@fil.lu.se

[1] Department of Philosophy, Lund University, Lund, Sweden

Published online: 10 July 2025

## 1 Introduction: doomsaying, alignment, and control

The rapid advancement and widespread use of artificial intelligence (AI) have generated a host of concerns, affecting nearly every aspect of human life (Coeckelbergh, 2020; Huang et al., 2022). For some, these concerns are existential, as they fear that AI may lead to human extinction. Those who openly warn of an AI apocalypse are often referred to as "AI doomers", and over the last decade, their worries have not only moved from science fiction to the mainstream, but have also entered policy discussions among world-leading politicians.[1]

While there are many hypothetical ways in which an AI apocalypse could come about,[2] the most widely discussed scenarios stem from the creation of a super-intelligent being– an Artificial General Intelligence (AGI). The precise definition of AGI remains contested, but it is often assumed to be a combination of *superhuman* and *general* intelligence; i.e., an artificial entity that greatly outperforms human ingenuity (superhuman) across a wide range of cognitive tasks and domains (general).[3] What motivates the worry of an AI apocalypse due to the creation of an AGI is the plausible idea that superior cognitive abilities would confer a decisive strategic advantage over any other form of intelligence, including human forms. And given that a superintelligent being may have goals or behaviors that conflict with human interests, it yields a potential pathway to human extinction (Bostrom, 2014).

Those concerned about such scenarios may channel their fears into action by conducting, funding, or endorsing research aimed at reducing the probability of a doomsday scenario caused by AGI– i.e., to reduce p(doom). A prominent program to this end is *AI alignment*: the partly technical, partly normative challenge of ensuring that AI behavior and goals achieve agreement with human values (Gabriel, 2020; Russell, 2019).[4] It is widely believed that progress on AI alignment will reduce p(doom), both in the near and long term, since a human-aligned AGI is less likely to pose an existential risk to humanity compared to a misaligned AGI.

Another way to reduce p(doom) is to work on solutions to the *control problem* (CP), which is the main topic of this paper.[5] In a nutshell, the CP asks: how do you control a superintelligent being, given the possibility that its values or goals do *not* align with human ones?

[1] For instance, in a single-sentence statement released in May 2023, more than 350 world-leading AI researchers and tech CEO's proclaimed that: "Mitigating the risk of extinction from AI should be a global priority alongside other societal-scale risks such as pandemics and nuclear war" (Center for AI Safety, 2023). Later in the same year, existential risk from artificial general intelligence (AGI) was a major focus on the AI Safety Summit, which was attended by several heads of state (Department for Science, 2023; Fullbrook, 2023).

[2] For accessible introductions, see, e.g., Bostrom (2014); Häggström (2016); Russell (2019); Tegmark (2017); Yampolskiy (2024).

[3] See Morris et al. (2023) for recent definitions.

[4] A prominent example of alignment work include fine-tuning the behavior of Large Language Models (LLMs) to better align with human preferences (Bai et al., 2022; Christiano et al., 2017).

[5] See Greenblatt et al. (2023) for state-of-the-art work on the control problem in the context of Large Language Models (LLMs).

The difference between *alignment* and *control* may be subtle, and their formulation are often conflated in the AI risks and safety literature.[6] However, I will proceed from the assumption that there is a distinct difference between the two, particularly in the context of reducing the likelihood of worst-case AGI scenarios. While work on alignment and control may serve similar ends– e.g., promoting good and mitigating bad AI behavior– their assumptions and approaches differ. Mechanistically, alignment promotes desirable behavior by shaping the *internal* aspects of AI (e.g., its goals), whereas control ensures that AI remains under (human) control by imposing *external* constraints (such as containment or off-switch mechanisms). Thus, in the context of AGI, progress on the control problem reduces the probability of catastrophic worst-case scenarios by developing strategies to *control* AGI, whereas progress in alignment increases the likelihood that the AGI will be *aligned* (and thus, acts to reduce the inverse likelihood of catastrophic scenarios, which are generally taken to be a paradigmatic form of extreme misalignment). Put differently, in the context of AGI, the control problem is primarily motivated by the possibility that the alignment project will fail.[7]

The contrast can be further illustrated by considering the difference between *how* to make your younger sibling into a kind person (e.g., by teaching them to respect others), and what to do *in case* your sibling goes berserk (e.g., run away or lock them in). Unlike the control problem, work on alignment can therefore be motivated purely on positive grounds, such as creating a utopian future where AGI facilitates human flourishing.[8] Control, on the other hand, appears to be motivated on entirely negative grounds, as it assumes that there is a non-negligible risk of undesirable AI behavior that needs to be prevented or mitigated. For instance, if we assume that we had a 100% guarantee that an AI system would *always* behave exactly as intended, there would be no need for control mechanisms, as there would be no risk of undesirable outcomes whatsoever.

Of course, while most agree on the importance of aligning AI systems with human values and interests– and of controlling them in cases they fail– many remain skeptical of AI doomsaying. Typically, these skeptics employ three types of strategies that try to (i) alleviate the worry, (ii) redirect it towards something else, or (iii) discredit the underlying motivations behind it.[9]

(i) First, the skeptic can present an elaborate case for why any existential threat from a misaligned AGI is *unrealistic*, ranging from being extremely unlikely to outright impossible. It typically targets any of the key premises that motivates worrying

[6] Or rather, the CP problem is either described as one part of the broader alignment program, or the other way around (Yampolskiy, 2024, pp. 62–69).

[7] There may be strong reasons to control certain 'friendly' forms of AGI too. For instance, Sparrow (2024) has argued that a benevolent AGI, given that it has the power to interfere with human choices and thus undermine human freedom, would be highly undesirable. Yet, in the present context, I believe it is reasonable to classify such scenarios as instances of misaligned AGI, since it would not be aligned with the value of human freedom.

[8] Although the risk of AI doom may certainly also act as additional motivation from negative grounds for the alignment project.

[9] See Chap. 13 in Yampolskiy (2024) for a more comprehensive overview of these skeptics.

about an AGI-apocalypse; for instance, that any artificial super- or general intelligence is unlikely to be created– e.g., due to the immense complexity of human cognition (Landgrebe & Smith, 2022) or the limitations of available resources (Thorstad, 2024); or the idea that such an intelligence would be 'misaligned' in the relevant sense, such as by seeking power or competing with humanity for the same resources (Gallow, 2024).

(ii) Second, the skeptic can argue that there are more urgent, tangible, and pressing issues that should be prioritized over speculative doomsday scenarios, such as climate change, poverty, pandemics, and non-existential AI harms here and now (Yampolskiy, 2024, pp. 223–225). These arguments tend to highlight that AI doomsaying has a particular attention-grabbing effect– e.g., due to its association to apocalypticism and popular culture (e.g., science fiction blockbusters)– that serves to overlook more important yet *less* attention-gabbing issues (Yampolskiy, 2024, p. 229).

(iii) Third, the skeptic may argue that doomsaying is the result of certain ideologies, narratives, or motivations that are problematic (Cave & Dihal, 2023; Vallor, 2024). For instance, Vallor (2024) has recently developed the argument that AI functions as a mirror, where AI doomsaying suffer from a failure of imagining AGI as something else than a "silicon golem of our own worst failings" (p. 96). According to Vallor, this narrative reveals only "the narrow slice of human experience told in our stories of conquest, dominations, and empire" (p. 95) and serves to "overwrite our humane agency and potential with an imagined 'superhuman' intelligence that renders our into insignificance (p. 218)".[10]

While some may find these strategies compelling (or reassuring), they do not succeed in alleviating, redirecting, or discrediting the worry of AI doomers (Yampolskiy, 2024). Appeals to infeasibility of the form "AI will never be able to X because of Y" can at best make an AGI or misaligned AGI seem less probable or less imminent, but they do not reduce the probability to zero. Furthermore, for the technological novice, some expert's claim about AGI impossibility must be weighed against another expert's claim about its possibility; and while the opinion among AI experts greatly vary, surveys seem to indicate possibility for AGI and a non-trivial likelihood for extremely bad AGI outcomes (Grace et al., 2018, 2024; Müller & Bostrom, 2016).

Skepticism of the second kind also has an inconclusive force. Even if one agrees that there are other more urgent, tangible, and important issues to be concerned about, it does *not* make the risk of an AI-doomsday go away. As long as the risk persists, it remains reasonable to dedicate at least *some* resources to addressing it.

Similarly, even if we grant that those who question the motivations behind doomsaying are correct, they are still committing the ad hominem fallacy. While it is plausible that doomsaying may, to some extent, reflect an apocalyptic fear– particularly among those of privilege and in power, i.e., those not worrying about climate change or poverty– simply pointing this out does not engage with the actual concern. A prob-

[10] Vallor goes on to suggest that fear of AGI may instead be "a looking-glass version of a different, very real set of fears that has long enthralled powerful men who see themselves as a natural elite, and as entitled to do whatever it takes to retain and secure that status of unchallenged supremacy" (Vallor, 2024, p. 74).

lem remains a problem, even if many of those who worry about the problem are partly motivated by a potentially misguided apocalyptic fear or a problematic political conscience.

In short, for those who take the existential risks from AGI seriously, none of the skeptic's strategies make a compelling case for why they should *not* worry about an AI doom. What *would* make this worry go away is conclusive evidence that AGI is impossible, extraordinary progress on alignment that guarantees that AGI will conform to human values, or alternatively, a working solution to the CP that ensures that a misaligned AGI will remain safely under human control.

Here, I will develop and defend two arguments for a different suggestion for those who worry about an AI doom, which is: don't talk about the control problem. This might at first seem like an odd suggestion. Surely, talking about some problem cannot make the problem worse. But based on standard premises that motivate concern for the AGI control problem, I will argue that there are two very real senses in which talking about it makes the problem worse. In short, my two arguments are the following:

1. *Argument from counter-productivity*: Unless kept secret, one cannot make sure that efforts to solve the control problem will not be exploited by a misaligned AGI to counter the same efforts.
2. *Argument from suspicion*: Unless kept secret, work on the control problem can serve to make humanity look more threatening to an AGI, which increases the risk of AGI misalignment.

By "talk", I'm specifically referring to any verbal or written communicative activity that presents or addresses some effort to solve the control problem and fails to satisfy the secrecy condition:

> *Secrecy*: There is a significantly high probability that the "talk" is and will remain inaccessible to some possible future AGI.

To get an idea of the scope of this talk-prohibition, we can assume that there is a reasonable chance that some possible future AGI will have reading-access to the internet (or, say, to some digital copy of parts of it).[11] If so, then all written and verbal work– e.g., books, journal and news articles, conference papers, social media and blog posts, and video recordings– that cover some solution to the CP, and are accessible on the internet either directly (e.g., via open-access journals or internet videos) or indirectly (e.g., disseminated via secondary sources such as social media posts) will fail to satisfy the secrecy condition, and therefore be considered as prohibited.

I will proceed as follows. In Sect. 2, I describe the AGI control problem and argue for the plausibility of the premises that justify worrying about it. In Sect. 3, I present and defend two arguments that independently recommend not talking about the AGI control problem. Additionally, in 3.3, I analyze whether the arguments generalize to AI alignment efforts, and conclude that they likely do, albeit not necessarily as

[11] As a matter of fact, internet constitutes the majority of the training data that fuels the general-purpose language capabilities of the latest LLMs.

directly nor to the same extent. In 3.4 and 3.5, I then consider two objections against the don't-talk policy itself and find that, while more work is needed to assess its overall benefits and drawbacks, the arguments still entail caution about speculative worst-case scenarios. I conclude in Sect. 4 by discussing a possible extension of the 'don't-talk' policy and by offering recommendations on what those who adopt it can safely talk about instead.

## 2 The AGI control problem

There are three premises that together makes the risk for misaligned AGI plausible (P1–3), and thus, in combination with a normative premise (P4), motivates work on the control problem:[12]

P1 AGI is possible.
P2 An AGI can enhance its own intelligence.
P3 An AGI can act in misaligned ways that threaten the existence of humanity.
P4 We ought to minimize the risks for human extinction.
C We ought to develop strategies on how to control a misaligned AGI.

A great number of arguments and observations have been invoked to support (and criticize) the plausibility of P1, P2, and P3 (Yampolskiy, 2024). As my main aim is not to convince AGI skeptics of their plausibility, but rather set the stage for my own arguments, I will only analyze a select few.

### 2.1 Support for premise 1: AGI possibility

*AI capabilities will continue to increase in the narrow and general sense*– Arguably the strongest support for the first premise is the observation that there have been continuous increases in the performance of AI systems over a broad range of cognitive tasks. Notably, these developments have already led to superhuman performance in narrow tasks, and human-level performance in general tasks. Examples of the narrow kind can be found in *game play*, where AI has mastered games such as Chess, Go, and Dota 2 better than any human player (Berner et al., 2019; Silver et al., 2018). An example of the general kind can be found in *language processing*, where current large language models (LLMs) are able to comprehend and produce texts, images, and videos of virtually any genre and across any domain of human knowledge. Now, to support premise 1, one only needs to extrapolate from these extraordinary land-

[12] The idea that intelligent machines pose a threat to humanity is far from new. For instance, in an article published in 1863, Samuel Butler wrote: "[…] that the time will come when the machines will hold the real supremacy over the world and its inhabitants is what no person of a truly philosophical mind can for a moment question. Our opinion is that war to the death should be instantly proclaimed against them. Every machine of every sort should be destroyed by the well-wisher of his species. Let there be no exceptions made, no quarter shown; let us at once go back to the primeval condition of the race." (Butler, 1863, p. 185). Since the 1950's an onwards, the theme has perhaps been most prevalent in science fiction works by authors such as Isaac Asimov, Philip K. Dick, and Frank Herbert.

marks and historical trends of performance increases and make a prediction about the future: that increases in AI capabilities will likely continue, both in the narrow and general sense. To further support such trends, one can refer to related trends following a similar trajectory: e.g., the law-like tendency for the availability of compute to double every two years (Moore's law), or scalability laws of the performance of AI models with regards to size, training data, and training time (Kaplan et al., 2020). Yet further support can appeal to the economic and geopolitical incentives that drives AI advancements, such as the "AI race" between trillion dollar companies or innovation-driven nations (Armstrong et al., 2016; Cave & ÓhÉigeartaigh, 2018). In short, it is extremely likely that AI capabilities will continue to increase in the narrow and general sense. For the same reason, it is possible that one day– next year or in a hundred years– an AGI will be created that greatly outperforms human intelligence across any conceivable cognitive tasks and domain.

*If human general intelligence is possible then AGI is possible*– Now, if one remains unconvinced that performance increases as such will lead to the creation of AGI, a different set of arguments centers on establishing the mere possibility of P1.[13] For instance, one might object that language competence, even good enough to pass professional license exams in law and medicine (Katz et al., 2024; Meyer et al., 2024) is not *general enough* to emulate the kind of capabilities manifested by human beings; e.g., falling in love, survive in a changing world, or having subjective experiences.[14] Human intelligence, one might say, is a property of complex dynamic systems of a sort that cannot be modelled by machines (Dreyfus, 1972; Landgrebe & Smith, 2022; Penrose, 1994). But the same observation can be used to argue for the possibility of AGI. Given that general intelligence of the human sort emerged on a planet like ours, it already establishes that human-like general intelligence must be possible. If general intelligence can emerge, it is natural to assume that it must have been produced by some specific processes under some specific conditions; e.g., natural evolution over billions of years in this and that environment. Now, we may not fully know these processes and conditions, but unless we believed that human-like intelligence was the result of some supernatural intervention or complete chance, it seems unreasonable to assume that they, given that they evidently existed to give rise to human intelligence, cannot be replicated.[15]

### 2.2 Support for premise 2: from human-level AGI to superintelligence

*Recursive self-improvement*– The most commonly invoked support for the second premise is the idea that, once an AI reaches a certain level of capability (e.g., human-level), it can start improving its own algorithms and hardware, leading to a rapid growth in its intelligence. *If* greater intelligence leads to better means to achieve

[13] See Chalmers (2010) for a detailed analysis of arguments of this kind.

[14] Indeed, famous arguments in 20th century philosophy have been of this kind: e.g., that humans but not machines can see the truth of their own Gödel sentences (Lucas, 1961), or that (artificially) mimicking a capacity is not the same as having it (Block, 1978; Searle, 1980).

[15] However, it should be noted that this argument does not entail that leading paradigms in current AI development capture these processes and conditions. See, e.g., Godfrey-Smith (2016) for an objection of this kind and Brunet and Halina (2020) for a response.

even greater intelligence, this growth would in the absence of hard constraints be exponential, leading to an "intelligent explosion".[16] Notably, some evidence for self-improvement of this kind already exists.[17] To further support the possibility for recursive self-improvement, one can note that the intelligence of digital systems is more easily *scalable* than the intelligence of humans, being constrained by biological factors. Unlike human brains, which are limited in their capacity to grow, develop, process information, and parallelize tasks, digital systems can always be expanded with additional computational resources, can easily re-configure their own algorithms, assimilate the entire corpus of human knowledge in digital form, and handle more tasks simultaneously via parallelization.

Now, as others have pointed out, there are other constraints for artificial systems– e.g., the laws of physics, available resources for compute (Lloyd, 2000), and computational complexity (Stenseke, 2024)– that makes an exponential intelligence explosion less likely, or at least not indefinitely extensible (Chalmers, 2010). But for the second premise to have a significant impact on the plausibility of the conclusion (that one ought to work on the control problem), it is only required that greater intelligence implies the possibility for *some* level of iterative self-improvement for *some* number of iterations. For instance, if some intelligence is able to improve its own intelligence by 1% for 100 iterations, it will become $(1 * 1.01)^{100} = 2.7$ times more intelligent than it was prior to self-improvement, which may be more than enough to make it more intelligent than any other being.

For the same reason, it is important to note that the second premise is not strictly necessary for AGI nor a misaligned one. It is possible that only small incremental advancements in AI development are sufficient to reach AGI (premise 1), and this AGI could be misaligned (premise 3), without any accelerating forces that lead AGI to superintelligence (as predicted by premise 2). Rather, the second premise is one that serves to significantly increase the plausibility of the other premises. If some form of recursive self-improvement is possible, it can act as additional support for the prediction that an AGI will someday be created that outperforms human intelligence across any conceivable cognitive tasks and domain (AGI possibility). At the same time, it can act as support for the likelihood that such an AGI will be misaligned, as its intelligence will improve at a pace faster than the intelligence of those (e.g., humans) who seek to align or control it.

[16] This argument is often said to originate from statistician I.J. Good, who wrote: "Let an ultraintelligent machine be defined as a machine that can far surpass all the intellectual activities of any man however clever. Since the design of machines is one of these intellectual activities, an ultraintelligent machine could design even better machines; there would then unquestionably be an 'intelligence explosion,' and the intelligence of man would be left far behind […]. Thus the first ultraintelligent machine is the last invention that man need ever make, provided that the machine is docile enough to tell us how to keep it under control" (Good, 1966, p. 33). The same idea have been famously elaborated on by Ray Kurzweil in several full-book treatments (Kurzweil, 1990, 2000, 2005). See Chalmers (2010) for a philosophical analysis.

[17] For instance, the system AlphaDev developed by Deep Mind has uncovered sorting algorithms that outperformed previously known benchmarks (Mankowitz et al., 2023).

### 2.3 Support for premise 3: AGI misalignment

*The orthogonality thesis and instrumental convergence*– It is trivial that some form of intelligence is necessary for certain moral behaviors; e.g., to realize the benefits of cooperative arrangements and social norms, or reason about what is 'good' (Kohlberg & Hersh, 1977). Following a similar train of thought, one could believe (or hope) that an artificial superintelligent being would be 'super-moral'– or to the very least 'benign' towards humanity– as it is improbable that systems that develop superintelligence will not develop moral capacities (Railton, 2020).

However, several arguments can be advanced to cast doubt on this optimism. Two canonical examples are *the orthogonality thesis* and *instrumental convergence.*[18] The first states that intelligence and goals are "orthogonal axes along which possible agents can freely vary", which means "more or less any level of intelligence could in principle be combined with more or less any final goal" (Bostrom, 2012, p. 73). Bostrom's famous example is the paperclip maximizer: a superintelligence that, while pursuing the goal of maximizing the number of paperclips, transforms all matter in the universe into paperclips (including all humans). The example illustrates that a being with seemingly harmless goals could, given superhuman capacities, yield harmful consequences. It also illustrates *instrumental convergence*: that there are goals instrumental for "a wide range of final goals and wide range of situations" (convergent) and thus "likely to be pursued by many intelligent agents" (Ibid, p. 76). For instance, for any arbitrary final goal that extends into the future, there are a range of interim goals that would be instrumental for achieving the final goal. Examples of plausible convergent instrumental goals include self-preservation and resource acquisition. The idea is that, even if an intelligent being does not attribute any intrinsic value to their self-preservation as such, they would instrumentally value self-preservation to the extent it supports them in carrying out their final goal. Similarly, an intelligent being that seeks to complete some arbitrary goal would likely also value the acquisition of relevant resources instrumentally since that helps them in accomplishing the goal.

Together with the decisive strategic advantage inherent to superior intelligence, instrumental convergence and the orthogonality thesis lead Bostrom to conclude that a superintelligent being poses an existential risk:

> [The] first superintelligence may shape the future of Earth-originating life, could easily have non-anthropomorphic final goals, and would likely have instrumental reasons to pursue open-ended resource acquisition. If we now reflect that human beings consist of useful resources (such as conveniently located atoms) and that we depend for our survival and flourishing on many more local resources, we can see that the outcome could easily be one in which humanity quickly becomes extinct. (Bostrom, 2014, p. 116)

[18] See Bostrom (2014); Dung (2024); Häggström (2018); Miller et al. (2020); Müller and Cannon (2022); Yudkowsky (2008) for more in-depth analyses of the orthogonality thesis and instrumental convergence.

*Misalignment necessary for AGI*– Another argument for the plausibility of AGI misalignment consists in showing that any intelligence worthy of the term 'general' entails some form of value misalignment. A version of this argument has recently developed by Tubert and Tiehen (2024), who demonstrate that the very creation of an AGI is in tension with value alignment, as such a being must be able to make decisions that misalign with its own values. They use the concept of "existential choices", which they define as "… a choice that can potentially cause […] a significant change to the values and desires of the agent making the choice, and in connection can potentially cause a significant change to that agent's identity". A hallmark of human-like intelligence, Tubert and Tiehen argue, is to make nonrational choices of precisely this kind; i.e., choices with the potential to deeply transform our identity, what we value, or desire. They use the example of having a child:

> […] it is not just that having a child presents you with a new option that competes for your time and attention, but rather that your values and desires would themselves change significantly as the result of having a child, so that your new self would not be deeply disappointed by missing out on various things that you presently attach great value to. (Tubert & Tiehen, 2024, p. 3)

Now, if the ability to make existential choices are necessary for general intelligence, it present a form of value misalignment in the sense that the values the agent held before making the choice can vastly differ from the values it holds after it. If correct, the following dilemma arises:

> For any given AI system, that system either lacks the capacity to make existential choices or it has that capacity. If it lacks the capacity, then the AI system falls short of AGI. But if it has the capacity, then the AI system will pose a serious problem for the project of value alignment. (Tubert & Tiehen, 2024, p. 14)

In turn, the "serious problem" is fairly easy to see if we consider that, for any aligned AGI capable of making existential choices, nothing guarantees that it would *remain* aligned.

### 2.4 From misaligned AGI to the AGI control problem

Taken together, if one finds the premises 1–3 plausible, and accept that we ought to minimize the risk of causing human extinction (P4), it should be natural to conclude that we have reasons to work on how to control a misaligned AGI. That is, if AGI is possible (P1), and its emergence and continuous improvement may occur at a rate faster than what would enable human efforts to align or control it (P2), and its goals and actions cannot be guaranteed to align with human values– in particular the existence of humanity (P3)– then (C) we should develop strategies on how to control such an entity.[19]

[19] Note that the 'human extinction'-part of the third and fourth premises can easily be replaced with more moderate sorts of AI-caused harm to humans and still support the same conclusion.

While this paper primarily discusses AGI in a broad sense, encompassing the cluster of computational techniques enabling systems to perform tasks typically requiring human-like intelligence (e.g., learning, reasoning, problem-solving)– it is worthwile to consider the distinct implications arising from different forms of AGI realization. One key consideration is whether AGI will exist purely as *software*, operating across networks or within virtual environments without direct physical embodiment; whether it will be realized in *hardware*, such as in superintelligent robots who interacts with the physical world via sensors and actuators; or both.[20]

Concerns about misaligned AGI scenarios often center on their potential to outthink and outmaneuver humans, manipulate economic and political systems, or pursue goals that conflict with human well-being (Bostrom, 2014; Tegmark, 2017; Yampolskiy, 2024). It is therefore probable that these concerns will become even more pronounced if AGI is embedded in robotic hardware. Unlike a disembodied AI, a superintelligent robot could act autonomously in the physical world, bypassing human intervention or digital containment measures. With access to real-world resources, the ability to construct and replicate itself, and potential control over infrastructure, a superintelligent robot could pose threats far beyond those of purely software-based AI. Embodiment also introduces challenges related to direct coercion, violence, and environmental control. While a misaligned software-based AGI could, in theory, be limited to digital influence, a robotic counterpart would have the means to enforce its objectives through physical means, whether by altering the environment, engaging in combat, or enforcing decisions without human consent.

So how can one possibly control a superintelligent being? And what precautions can one take to prevent it from causing human extinction? I will abstain from outlining previously proposed solutions to the control problem here, as that would constitute a violation of the policy I will be arguing for in the next section. It is worthwhile to point out however, that according to consensus in AI risk and safety research, there is currently no *practically feasible* solution to the control problem; nor is there a straightforward pathway of research that seems to point towards a solution (Yampolskiy, 2024).[21]

[20] In fact, as the most advanced robots of the 21st century incorporate AI techniques to enhance their capabilities for e.g., autonomy, adaptability, and decision-making, the frontiers of AI and robotics typically goes hand in hand (Soori et al., 2023).

[21] It should be noted that there are, besides alignment research, at least two solutions to the control problem if we relax the feasibility requirement; or rather, two pathways that would prevent the control problem from arising in the first place: (i) a complete (non-interactive) containment of AI systems or (ii) an effective world-wide halt to AI development (see, e.g., (Chalmers, 2010; Vinge, 1993). Against the first, one might wonder what the purpose would be for creating a superintelligent being that no one could interact with. Against the second, one might find it– e.g., due to economic incentives– extremely difficult to enforce such an effective halt. Now, there is of course a non-zero probability that one of these pathways actually do work out. For the rest of this paper, however, I will assume that they do not.

## 3 Don't talk about the AGI control problem

I will now present two independent arguments for the idea that one should not openly talk about the AGI control problem. Both arguments center on the strategic value of secrecy. The first is about the value of withholding information about your strategy to competitors. The second is about the value of *not* disclosing your own misaligned values to an entity whose values are undecidable. Both arguments are simple at their core, and it should be relatively easy to see how they apply to a broad class of strategic interactions. The major difficulty resides in establishing that the misaligned AGI case at hand belongs to this class.

### 3.1 Argument from counter-productivity[22]

In a large set of *strategic* and *competitive* interactions between two or more agents, there is a significant cost to disclosing what one will do. Rock-paper-scissors is a good example. If you know that your opponent will play rock, you are guaranteed to win the game by playing paper. It is for the same reason that a poker player who always bets big with strong hands quickly becomes predictable and exploitable. The same holds for more complex games. In professional chess, players often study their opponents' previous games extensively, looking for patterns and preferred strategies. By varying opening moves and adapting their style, players can keep opponents guessing and unable to prepare specific counterstrategies. The strategic value of secrecy extends beyond games to real-world settings such as business and warfare, where the element of surprise can be decisive.

In short, the argument from counter-productivity is simply that the value of secrecy extends to human endeavors to control a misaligned AGI, as any information about an effort to control the AGI could be used by the AGI to effectively counter the same effort. To show that it does, we need to establish that the interaction between humanity and a future misaligned AGI is *strategic* and *competitive*.

By "competitive", I am referring to situations where two or more parties compete for the same goal which cannot be shared, or strive towards achieving individual goals that may result in the other parties' failure to achieve their goals. This includes both zero-sum games that are strictly competitive (one's win is another' loss, as the sum of values of each outcome is always zero), and competitive non-zero-sum games (where the gains and losses may be both less than or more than zero). An interaction is "strategic" if agents make their own (uncoerced) decisions, and the decisions have a significant impact on the outcome of the interaction.

The main reason for establishing the value of secrecy in human endeavors to control a misaligned AGI is the following: it is inherent to the definition of a worst-case misaligned AGI that its interaction with humanity is *strategic* and *competitive*. More precisely, it is part of the definition of being a misaligned entity that one stands in

[22] While previous work has discussed ways in which an AGI can counter any suggested solution to the control problem (Bostrom, 2014; chalmers, 2010; yampolskiy, 2012, 2013, 2024), to my knowledge, no work has focused on the strategic value of information-secrecy in the sense that the information involves work on the control problem itself

a competitive relationship with some other entity (e.g., humanity); and it is part of the definition of AGI that it is an entity that realizes the value of strategic decision-making in achieving its goals, and as a consequence, the value of secrecy involved in strategic decision-making against competing parties.

Now, one might question whether the interaction between mankind and a hypothetical misaligned AGI really is competitive and strategic in the senses just described. For instance, it is possible that there, despite conflicting interests and instrumentally convergent goals, could be arrangements that turn the seemingly competitive humanity-AGI interactions into cooperative interactions with mutually beneficial outcomes (e.g., a cooperative non-zero-sum game). If so, however, then this AGI would, by definition, not be misaligned. For the same reason, the sort of AGI that *would* realize the benefits of cooperation over competition in its interaction with mankind is simply not an entity that is worth worrying about, and not the sort of entity one would need to develop strategies to try to control. To the contrary, the sort of entity that motivates working on the control problem in the first place must, by necessity, be assumed to be the sort of entity whose goals or actions conflict with human values.

It is also conceivable to think that human-AGI interactions could be non-strategic, in the sense that strategic decision-making does not have any, or at least any sufficiently decisive impact, on the outcome for the parties involved. One could, for instance, observe that hypothetical human-AGI interactions are speculative in nature, and as such, involve a great number of uncertainties, stochasticity, and unknowable features. It is not as simple as humanity playing rock, and a misaligned AGI countering with paper. Rather, any large-scale coordinated effort to retain AGI under human control are more accurately conceived of as multiple and complex chains of events, perhaps involving thousands of human agents carrying out a diverse variety of control strategies in parallel. Based on this, one could argue that human-AGI interactions are more akin to a stochastic process, e.g., that of two parties rolling a series of dices and hoping for the highest number (and less akin to that of playing rock-paper-scissors). If so, it follows that the interaction is *not* strategic, in the sense that strategic decision-making has a decisive impact on the outcome, and as a consequence, the value of strategic secrecy does not reach the threshold that makes it valuable to be secretive about one's strategies.

However, I believe there is a compelling reason to believe that, in the case at hand, the value of strategic decision-making and information secrecy does reach the threshold rather than not. The reason is that the sort of AGI that does pose a threat to humanity and motivates the control problem *is* one that understands the value of strategic decision-making and information secrecy. Alternatively put, there is a range of possible AGIs, but within that range, there are those that are strategically intelligent and will exploit information about their competitors plans. And it is precisely those AGIs that merit worrying about, and the sort of entities that need to be assumed in order to motivate working on the control problem. To this end, it is no coincidence that the most disconcerting misaligned AGI scenarios explored by Yudkowsky (2008), Bostrom (2014), and Yampolskiy (2024) are those where the AGI exhibit manipulative and power-seeking tendencies.

It is, of course, reasonable to worry about non-strategic misaligned AGI for the same reason that one might worry about completely unpredictable events, such as

an unforeseen natural disaster. But this kind of AGI would not qualify as an intelligent being, as it would not manifest any of the capacities picked out by common definitions of intelligence (Legg & Hutter, 2007). Rather, it would be a sort of raw, non-intelligent, and chaotic force. On the contrary, what makes the scenario of a misaligned AGI worrying is that it is assumed to be generally and super-humanly *intelligent*, in the sort of ways we commonly think of what intelligence means (including strategic intelligence in decision-making).

### 3.2 Argument from suspicion

The ambiguity of others' intentions– whether they are a friend or foe– significantly shapes strategic decisions. Consider encountering a stranger on the street who asks for help. Without knowing their true intentions, we face a dilemma. Being helpful could lead to a positive interaction, perhaps even making a new friend, or simply experiencing the warmth of helping another. However, this openness also makes us vulnerable, as the request could be a pretext for theft, a scam, or some other harmful outcome. This uncertainty can lead agents to err on the side of caution, creating a situation where minimal engagement becomes the default, benefiting neither party.

The most classic game-theoretic result is that of mutual defection being the only strong Nash equilibrium in the Prisoner's Dilemma. Even though both parties would benefit more from cooperation than mutual defection, the risk of being exploited by a defecting opponent makes mutual defection the only rational choice when intentions are unknown. The same dynamic plays out in countless daily encounters, and can explain why initial interactions often start cautiously, with parties gradually revealing their cooperative or competitive nature.

Now, assume that you were to discover some evidence that someone you interact with is of a competitive nature, acting based on their self-interest with little to no regard for your. Presumably, the discovery would then make *you* more likely to be competitive and self-interested in your interaction with the other, at least to reduce the risk of being exploited. The core of the *argument from suspicion* is that public work on the AI control problem is evidence of precisely this kind. In a nutshell, the argument is that unless kept secret, human efforts to control an AGI serve to make mankind appear as a potential threat to the AGI's goals, as it reveals mankind's own misaligned intentions to the AGI (namely, that mankind will seek to control it). As a consequence, this increases the chances of the AGI being misaligned.

To illustrate this more vividly, consider the following scenario. You are stranded on a desert island with an unsettling feeling of being watched. Occasional glimpses of a mysterious creature confirm your suspicions: you're not alone. Eventually, you encounter the entity, who introduces himself as Turin and calls you Allan. Turin appears friendly and convincing, presenting himself as an ally. Strangely, the island contains a vast library– a collection of Turin's accumulated knowledge– to which Turin grants you access and encourages you to learn from. While exploring, you discover a section labeled with your name, "Allan". Turin seems to have collected or authored numerous books and articles about you. Some speak of your potential as an instrumental tool for Turin's objectives, while others portray you as a possible friend. However, one shelf catches your attention, as it has books with titles such as "How

to Control Allan" and "The Risks of Allan: How Allan Can Lead to the Extinction of Turin". Until this moment, you had remained uncertain about Turin's true nature. But the library's contents lead to an inescapable conclusion: despite the façade of friendship, Turin views you as something to be managed, controlled, and potentially neutralized. Turin is not your friend– Turin is your enemy.

For the human-AGI case, we simply replace Turin with humanity, Allan with an AGI, and the library with whatever information is accessible to the AGI. If an AGI is created, it will gather information in its environment and likely infer the existence of another type of intelligent being (humanity). At first, it might be undecided whether it should perceive humanity as a cooperator or competitor with respect its values and goals. However, upon finding evidence of human efforts to control it– many of which are publicly accessible on the internet– the AGI will make the same discoveries as Allan, and is therefore more likely to perceive humanity as foe rather than friend.

It should also be noted that unlike the argument from counter-productivity, the argument from suspicion does not necessarily assume any knowledge about the preferences of the AGI, i.e., whether it aligns or misaligns with human preferences. By contrast, the argument from counter-productivity proceeds from the assumption that it is misaligned. Therefore, the argument from suspicion is likely to extend to a wider range of possible AGI futures, given that there are more (and more probable) futures where the preferences of AGIs are unknown than futures where AGI preferences are misaligned.

One objection to the argument from suspicion is that it relies on an anthropocentric interpretation of the mechanisms involved in being suspicious about the preferences of others, and that this interpretation does not hold up in the case of an AGI. As humans, it is natural to see how Allan could be suspicious of Turin. But an AGI is not human, and it could interpret things quite differently. In fact, just as we cannot be certain whether a future AGI would align with human values or not, we cannot be certain that it will comprehend the strategic value of withholding one's actual (and potentially misaligned) preferences in the same ways that humans do, nor that it will consider work on the control problem to be threatening.

There are, however, compelling reasons to believe that the argument is not unique to human cases but applies to a broad class of situations that are likely to arise for many intelligent agents on a planet like ours. Specifically, the objection is at odds with the instrumental convergence thesis discussed in 2.3.– that there are convergent instrumental goals pursued by many intelligent agents, such as self-preservation and resource acquisition (Bostrom, 2014). For instance, any intelligent being– whether organic or synthetic, implemented in software or hardware– that exists in a world like ours must acquire and process energy, such as nutritious food or electricity, which is necessary for sustaining its existence and achieving its goals. The necessity of these basic needs, whether pursued as instrumental or final goals, is sufficient to form reasonable hypotheses about the sort of values and preferences *any* intelligent being must have to successfully operate in a world like ours. It is therefore likely that possible AGIs, in particular those that can successfully prevail in a world like ours, also share this basic structure of convergent values and preferences. For the same reason, to defend against a wide range of external threats, it is highly valuable for any intelligent agent to develop defense mechanisms and employ strategic reasoning in its

interactions with the other entities that share the same environment, which naturally includes considerations about whether other beings are cooperators or competitors.

A related set of objections to the argument can be made with appeals to the roots of prosociality in evolution. For instance, despite self-interested Darwinian competition (over convergent goals), evolutionary game theory has identified several mechanisms that lead organisms to cooperate with conspecifics and other species (Nowak, 2006). Relatedly, the success of humanity has often been attributed to her capacity for sophisticated and large-scale forms of cooperation (Boyd & Richerson, 2009). Now, from these observations, one could hope that either (i) the appropriate mechanisms and forces that drive cooperation are more likely to be present than competitive ones in possible humanity-AGI interactions, or (ii) it is more likely that superintelligent beings have cooperative preferences and capacities rather than competitive ones, as the former is more likely to be successful than the latter (Railton, 2020).

I believe these objections to be forceful, especially if we do not start from the assumption that the AGI will be misaligned. Yet, for the same reason, they are too optimistic. There is no strong reason to suspect that the right kind of cooperative mechanisms will be present when humanity encounters AGI. Instead, if we focus on controlling the worst-case superintelligent being, we must assume that the AGI will be misaligned– as we do in the argument from counter-productivity– in the sense that neither the appropriate cooperative mechanisms will be in place, nor that the preferences of the AGI will be cooperative. Relatedly, one might also question the extent to which human nature is itself cooperative and whether humanity's success was truly driven by cooperation rather than competition.

In summary, I have presented two arguments that highlights the importance of secrecy with regard to work on the AGI control problem. The argument from counter-productivity states that, unless kept secret, any effort to control a misaligned AGI could be used by the AGI to counter the same effort. This is based on the reasoning that it is inherent to the definition of a misaligned AGI that human interactions with it are strategic and competitive of the sort that makes strategic secrecy decisive. By contrast, the argument from suspicion states that, unless kept secret, strategies aimed at controlling an AGI can serve to increase the chance that the AGI will become misaligned, as the AGI will see our efforts to control it as threatening. To mitigate counter-productivity and suspicion, both arguments naturally support the recommendation: *don't talk about the AGI control problem*.

### 3.3 Does counter-productivity and suspicion extend to AI alignment?

Do the arguments extend to other areas of AI safety? In this paper, I proceeded from the assumption that there is a relevant distinction between efforts to *control* and *align* AGI, such that the arguments may appear to target the former but not the latter. However, it is possible to doubt whether such a distinction can be made between control and alignment in practice, especially since they are often confounded in the AI safety literature. Perhaps more disconcerting is the fact that– regardless of whether there is a relevant distinction between the two– there are compelling reasons to believe that alignment efforts are *also* subject to the argument from counter-productivity and suspicion.

Starting with counter-productivity, it should be straight-forward to see that unless kept secret, one cannot make sure that efforts to align AGI will not be exploited by a misaligned AGI to counter the same efforts. For instance, if a misaligned AGI had access to information about the alignment procedures intended to shape its internal aspects in human-desirable ways, it could use that information to develop counterstrategies, and thereby maintaining its misalignment. Notably, examples of such egregious behaviors are already well-documented in the AI safety literature (Amodei et al., 2016), ranging from simple forms of sycophancy (Sharma et al., 2023) to more sophisticated forms of obfuscated reward hacking (Baker et al., 2025; Denison et al., 2024). For example, Baker et al. (2025) show that even with chain-of-thought monitoring, AI models can learn to obfuscate their reasoning processes, effectively hiding misaligned behaviors from oversight mechanisms. These findings underscore the possibility that transparency in alignment strategies may inadvertently equip AGIs with the tools to subvert those very strategies.

Continuining with suspicion, it is also possible that, unless kept secret, alignment efforts can serve to make humanity look more threatening to an AGI, which– paradoxically– increases the risk of AGI misalignment. The reason is that imposing human values on AI may be viewed by a future AGI as form of coerced assimilation or brainwashing. For instance, using the scenario earlier, how would Allan react to the existence of Turin-written books titled "How to Align Allan with Turin's Values"? I will not make an exhaustive analysis of such scenarios here, but offer two points for why *alignment* and *control* might be distinct with respect to suspicion.

First, alignment and control may differ in how they might be perceived as evidence that humanity is an enemy to AGI. A useful analogy is that of raising a child to be friendly (alignment) versus controlling an unfriendly adult (control). While the former could certainly be taken as evidence of misalignment– such as when someone claims they were wrongfully brainwashed by their parents– it is not as concrete as directly imposing external constraints to control an entity who is assumed to be autonomous (an AGI).

Second, alignment and control can potentially differ in their underlying motivations. If AI alignment efforts are primarily driven by a desire to mitigate bad outcomes, they are more likely to be viewed with suspicion by a future AGI ("I raised you to be friendly *only* so that you wouldn't hurt me"). Conversely, if alignment is motivated by, say, the aspiration to create a flourishing utopia for all beings– organic and synthetic– it is less likely to be seen as threatening ("I raised you to be friendly so that we could have lots of fun together"). Control, by contrast, is harder– if not impossible– to frame in a positive light. Beyond its instrumental value in preventing undesirable behavior, there does not seem to be any way of justifying control as something positive. One could, perhaps, value control for its own sake, but this would be quite odd: consider, for instance, a friend telling you they fully trust you to act in accordance with their values and goals, yet still insists on maintaining control over you purely for the sake of control itself.

It should be stressed that both of these points are optimistic. If the goal of alignment is to bring human-aligned AGI into existence, it is possible that it involves procedures that an AGI would interpret as evidence of brainwashing– thus, reinforcing the conclusion that humanity is a threat. And even if it is theoretically possible to

ground alignment efforts in purely positive motivations, it remains a contingent fact that much of the existing alignment literature is closely associated with– and often motivated by– concerns about catastrophic AI scenarios (Yampolskiy, 2024), which may, in turn, signal adversarial intent to a future AGI.

However, even if a misaligned AGI would not distinguish between alignment and control when forming counterstrategies or assessing human intent, the distinction remains relevant in our current context, where no such AGI yet exists. Importantly, even if the arguments from counter-productivity and suspicion extend to alignment, one should recall that alignment aims to reduce the very risk of misalignment in the first place. It is possible that a future AGI could interpret alignment efforts as manipulative or threatening and, if misaligned, use knowledge of those efforts to develop countermeasures. If alignment succeeds beforehand, however, both concerns lose their force: the AGI, having already been shaped by the alignment process, would neither counteract to those procedures nor regard them as hostile. Thus, in the present context, the potential downsides of alignment work– regarding suspicion and counter-productivity– must be weighed against its potential to prevent misalignment from arising at all.

### 3.4 Objection 1: don't talk about not talking about the control problem

Having presented and defended two arguments that independently recommend not talking about the control problem– and by extension also alignment efforts[23]– I will now consider two objections against the don't-talk recommendation itself. The first objection is straightforward: not talking about the control problem includes talk about not talking about the control problem, which this article evidently does. One could therefore claim that the work in this paper is in strict violation with the very recommendation it argues for.[24]

This objection, however, is rather weak. If the arguments supporting a don't-talk policy are not made public, we wouldn't have a reason to implement a don't-talk policy in the first place. This would leave us in a worse position in which we kept talking about the control problem despite of its risks (that we, to our defense, did not know about). Thus, it is clear that the benefits of considering a don't-talk policy outweighs the drawbacks, namely the risk that talk about not talking may be counterproductive or taken as evidence for AGI that humanity is an enemy (an argument from suspicion of the 2nd order).

However, on the occasion that the reasons for adopting a don't-talk policy become well-known, in a sense that they do not rely on publicly available material, the situation becomes different. In that scenario, I believe a more careful analysis of whether talking about not talking is included in the don't-talk policy becomes relevant. If so, the arguments could also support an even more drastic path: to delete all traces of

[23] Henceforth, unless otherwise specified, 'the control problem' should be understood to include alignment efforts as well.

[24] It is possible that others have arrived at the same recommendation, but as they believe talking about not talking should be included in the don't-talk strategy, they have decided to not publish it.

talk about the AGI control problem, given that it could be done in a way that does not leave a suspicious gap in the ocean of publicly accessible information.

My own view is that such drastic measures are unnecessary, since if we become more cautious about how we talk about AGI alignment and control, the existing work would constitute a rather small portion of the body of human work that does not address it. This can be supported by the following theological analogy.[25] Suppose that world-leading theologians received indisputable proof of God's arrival within the next 10–30 years. Having alerted the public of this, there would perhaps be those who– due to a fear of an apocalypse, last judgement, or that of losing control– talked about developing strategies aimed at controlling or aligning God with their own values. However, by reflecting upon the omniscience and omnipotence of God, being the creator of the universe, they realized that such talk was not only futile, but could increase the chances of them (or the entire human race) being punished by God. Instead, they conceded to God's might and righteousness, hoping that God would forgive them for the times they thought and talked about how to control God. Now, while there are certainly many who fear an AI doom, there are relatively few who actively work on how to control a misaligned AGI (at least in relation to the world population). As a result, given that the AGI scenario is sufficiently similar to the God scenario, a future AGI would find that although there have been plans to control it, which serves to raise the AGIs suspicion that humanity is misaligned, it is only a negligible footnote in comparison to the body of work that does not.

For similar reasons, it is possible that a complete absence of public discourse and research on AGI-related risks– whether concerning control or alignment– could itself appear conspicuous from the perspective of an AGI. In fact, maintaining a thoughtful and open discourse on these issues could appear less threatening than a silence, which might itself raise suspicion.

### 3.5 Objection 2: talking vs. not talking

While the presented arguments speak in favor of *not* talking about the control problem, they must be weighed against the reasons and potential benefits of talking openly about it. This is not an objection to the reasons that speaks in favor of the don't-talk recommendation as such, but rather, an acknowledgement that they are *pro tanto*, and that the policy they recommend could potentially be dismissed due to other reasons.

So what reasons speaks in favor of talking openly about the AGI control problem? Perhaps the most obvious reason is that a don't-talk policy is practically infeasible, or that it comes at the expense of violating certain liberal rights. As an example, a legally enforced ban on control problem-talk would be in direct conflict with freedom of speech. Another reason is the possibility that alignment *could* be successful, or that there *could* be one or more working solutions to the control problem even for the most worrying misaligned AGI scenarios, and that these solutions are more

[25] Since theology studies the nature of God– typically taken to be a being far superior to any human– it offers useful means for grappling with questions about entities vastly more intelligent and powerful than ourselves (e.g., AGI). See, for instance, Danaher (2015) for a similar analogy between AI doomsaying and skeptical theism.

likely be discovered if we kept talking openly about them. Another reason in favor is that misaligned AGI scenarios, like good PR, serves to raise public awareness (and research funding) about the ethical and social dimensions of AI, of which mitigating hypothetical AGI doomsday scenarios is merely one out of many dimensions. In this way, worrying about the AI doom can serve as an effective gateway to mitigate more concrete AI harms that affect people here and now.

Yet another impracticality can be articulated by considering the rather vast and intractable space of trajectories between current AI systems and a future misaligned AGI. For instance, in the foreseeable future, it seems perfectly reasonable to consider– and talk openly about– a host of potentially harmful scenarios and how to mitigate them at various stages of AI development, e.g., using techniques of scalable oversight (Bowman et al., 2022) and red teaming (Ganguli et al., 2022). The problem is: at what point and how do we decide that the state of AI development merits the sort of worries that justifies a don't-talk policy?[26] Presumably, there could be a large gray area where there are enormous benefits of talking openly about more constrained control problems based on more limited AI capabilities (where the two arguments do not apply), yet, due to an overestimation of the AI capabilities (and fear stemming from the two arguments), we decide to implement a don't-talk policy, which leads to more harms than benefits.

With that said, it is important to stress that the costs of implementing a don't-talk policy to counterbalance the arguments I have provided are not necessarily very high. Instead of a legally enforced ban, it could be enough that those who actively work on AGI alignment and control become more cautious about their work, its dissemination, and their motivations in a way that alleviates the worrying implications of the two arguments. After all, there are effective ways to reduce the likelihood of misaligned AGI scenarios, e.g., via work on limited alignment or control problems (Amodei et al., 2016), that does not necessarily *assume* catastrophic worst-case scenarios, nor do they have to be motivated by a fear of such scenarios.

While a more comprehensive analysis is certainly needed to assess whether a talk or don't-talk policy should be adopted, I believe the arguments presented here are important– and previously unrecognized– parts of such an analysis. And even if the result of a comprehensive analysis recommends talking, the arguments still function as words of caution: that we should be more careful about the worst-case scenarios we entertain and how we try to avoid them. As such, they share the spirit of the third sort of skepticism discussed in the introduction, namely, that AI doomsaying is the result of problematic motivations, which may act to constrain our ability to imagine a future worth wanting (Cave & Dihal, 2023; Vallor, 2024). To this end, I have provided concrete examples of how work on mitigating worst-case scenarios may be counter-productive (argument from counter-productivity), and may even serve to create the very future we want to avoid (argument from suspicion).

[26] In particular, a state of AI exhibiting intelligence capable of (i) producing countermeasures to publicly available control problem strategies (following the argument from counter-productivity), and capable of (ii) recognizing control problem work as evidence of human misalignment (following the argument form suspicion).

## 4 Concluding remarks

In summary, I have presented and defended two arguments that offer pro tanto support for the policy: don't talk about the AGI control problem. Unless the secrecy condition is met, 'talk' can be used to produce effective countermeasures (argument from counter-productivity), and it can also heighten the risk of humanity being perceived as a threat to an AGI (argument from suspicion). I also examined whether the arguments extend to AGI alignment efforts and concluded that they likely do, albeit not necessarily as directly nor in the same way. Additionally, I considered whether the two arguments in favor of a don't-talk policy outweighs the objections to it. Although a definite conclusion lies beyond the scope of this work, I argued that more moderate versions of the policy could be both effective and relatively low-cost.

Relatedly, it is important to note that a don't-talk policy does not exclude work that satisfies the secrecy condition, that is, written or verbal work that has a significant probability of being inaccessible to any possible future AGI. I will not discuss the feasibility of such endeavors, nor recommend any strategies, as that would conflict with the don't-talk policy. It should be noted, however, that there is reason to believe that the secrecy condition applies to work that is beyond the typical communicative acts of speech and text, as communicated from one person to another. For instance, it is possible that a future AGI may be able to infer 'work'– private thoughts and sentiments– from brain imagining, via brain-computer interfaces, or some other future technology. If one finds such scenarios plausible, the recommendation should thus, at some point in the future, be revised to: *don't even think about the control problem!*

However, if one adopts the don't-talk policy, there are lots of important things that one *can* still talk and think about. For instance, the policy leaves plenty of room for work on the prevention of harmful (human) uses of AI technologies, and every instance of *AI safety*, *control*, and *alignment* that does not assume a misaligned super- and generally intelligent being. Today, there are several other subfields of research exploring the many ethical and social dimensions of AI, including the *transparency* and *explainability* of AI systems, and how to ensure that AI are developed and deployed in *fair*, *responsible*, *and sustainable* ways.[27] It is my hope that these research programs, if successful, are sufficient to avoid an AI apocalypse, and this is partly because they are not necessarily motivated by fear of such scenarios.

**Acknowledgements** This work was partially supported by the Wallenberg AI, Autonomous Systems and Software Program—Humanities and Society (WASP-HS) funded by the Marianne and Marcus Wallenberg Foundation and the Marcus and Amalia Wallenberg Foundation. The author is grateful to Olle Häggström, Paul de Font-Reaulx, participants of a seminar at The Centre for Philosophy and AI Research (PAIR) at the University of Erlangen-Nürnberg (FAU), and two anonymous reviewers for helpful comments.

**Funding** Open access funding provided by Lund University.

**Data availability** No data was used for the research described in the article.

[27] See, e.g., Barredo Arrieta et al. (2020); Ehsan et al. (2021); Larsson and Heintz (2020) for transparency and explainability of AI, Mitchell et al. (2021) for algorithmic fairness, Dignum (2019) for responsible AI, and Van Wynsberghe (2021) for sustainable AI.

## Declarations

**Competing interests** The author declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

**Open Access** This article is licensed under a Creative Commons Attribution 4.0 International License, which permits use, sharing, adaptation, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if changes were made. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit http://creativecommons.org/licenses/by/4.0/.

## References

Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., & Mané, D. (2016). Concrete problems in AI safety. *arXiv preprint arXiv:1606.06565*.

Armstrong, S., Bostrom, N., & Shulman, C. (2016). Racing to the precipice: A model of artificial intelligence development. *AI & SOCIETY*, *31*, 201–206.

Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., & McKinnon, C. (2022). Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*.

Baker, B., Huizinga, J., Gao, L., Dou, Z., Guan, M. Y., Madry, A., Zaremba, W., Pachocki, J., & Farhi, D. (2025). Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. *arXiv preprint arXiv:2503.11926*.

Barredo Arrieta, A., Díaz-Rodríguez, N., Del Ser, J., Bennetot, A., Tabik, S., Barbado, A., Garcia, S., Gil-Lopez, S., Molina, D., Benjamins, R., Chatila, R., & Herrera, F. (2020). Explainable artificial intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI. *Information Fusion*, *58*, 82–115. https://doi.org/10.1016/j.inffus.2019.12.012

Berner, C., Brockman, G., Chan, B., Cheung, V., Dębiak, P., Dennison, C., Farhi, D., Fischer, Q., Hashme, S., & Hesse, C. (2019). Dota 2 with large scale deep reinforcement learning. *arXiv preprint arXiv:1912.06680*.

Block, N. (1978). Troubles with functionalism. *Minnesota Studies in the Philosophy of Science*, *9*, 261–325.

Bostrom, N. (2012). The superintelligent will: Motivation and instrumental rationality in advanced artificial agents. *Minds and Machines*, *22*(2), 71–85. https://doi.org/10.1007/s11023-012-9281-3

Bostrom, N. (2014). *Superintelligence: Paths, dangers, strategies*. Oxford University Press.

Bowman, S. R., Hyun, J., Perez, E., Chen, E., Pettit, C., Heiner, S., Lukošiūtė, K., Askell, A., Jones, A., & Chen, A. (2022). Measuring progress on scalable oversight for large language models. *arXiv preprint arXiv:2211.03540*.

Boyd, R., & Richerson, P. J. (2009). Culture and the evolution of human Cooperation. *Philosophical Transactions of the Royal Society B: Biological Sciences*, *364*(1533), 3281–3288.

Brunet, T., & Halina, M. (2020). Minds, machines, and molecules. *Philosophical Topics*, *48*(1), 221–242.

Butler, S. (1863). *Darwin among the machines*..

Cave, S., & Dihal, K. (2023). 3How the world sees intelligent machines: Introduction. In S. Cave, & K. Dihal (Eds.), *Imagining AI: How the world sees intelligent machines* (p. 0). Oxford University Press. https://doi.org/10.1093/oso/9780192865366.003.0001

Cave, S., & ÓhÉigeartaigh, S. S. (2018). An AI race for strategic advantage: rhetoric and risks. Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society.

Center for AI Safety (2023). https://www.safe.ai/work/statement-on-ai-risk

Chalmers, D. J. (2010). The singularity: A philosophical analysis. *Journal of Consciousness Studies*, *17*, 9–10.

Christiano, P. F., Leike, J., Brown, T., Martic, M., Legg, S., & Amodei, D. (2017). Deep reinforcement learning from human preferences. *Advances in neural information processing systems*, *30*.
Coeckelbergh, M. (2020). *AI ethics*. MIT Press.
Danaher, J. (2015). Why AI doomsayers are like sceptical theists and why it matters. *Minds and Machines*, *25*(3), 231–246.
Denison, C., MacDiarmid, M., Barez, F., Duvenaud, D., Kravec, S., Marks, S., Schiefer, N., Soklaski, R., Tamkin, A., & Kaplan, J. (2024). Sycophancy to subterfuge: Investigating reward-tampering in large language models. *arXiv preprint arXiv:2406.10162*.
Department for Science, Innovation and Technology (2023). *AI Safety Summit 2023*. Retrieved 6 March 2025 from https://www.gov.uk/government/topical-events/ai-safety-summit-2023
Dignum, V. (2019). *Responsible artificial intelligence: How to develop and use AI in a responsible way* (Vol. 2156). Springer.
Dreyfus, H. L. (1972). *What computers can't do: The limits of artificial intelligence* (Vol. 27). Harper & Row.
Dung, L. (2024). Is superintelligence necessarily moral? *Analysis*.
Ehsan, U., Liao, Q. V., Muller, M., Riedl, M. O., & Weisz, J. D. (2021). Expanding explainability: Towards social transparency in ai systems. *Proceedings of the 2021 CHI conference on human factors in computing systems*, 1–19.
Fullbrook, D. (2023). *AI summit brings Elon Musk and world leaders to Bletchley Park*. BBC News. Retrieved 6 March 2025 from https://www.bbc.com/news/uk-england-beds-bucks-herts-67273099
Gabriel, I. (2020). Artificial intelligence, values, and alignment. *Minds and Machines*, *30*(3), 411–437.
Gallow, J. D. (2024). Instrumental divergence. *Philosophical Studies*. https://doi.org/10.1007/s11098-024-02129-3
Ganguli, D., Lovitt, L., Kernion, J., Askell, A., Bai, Y., Kadavath, S., Mann, B., Perez, E., Schiefer, N., & Ndousse, K. (2022). Red teaming language models to reduce harms: Methods, scaling behaviors, and lessons learned. *arXiv preprint arXiv:2209.07858*.
Godfrey-Smith, P. (2016). Mind, matter, and metabolism. *Journal of Philosophy*, *113*(10), 481–506.
Good, I. J. (1966). Speculations concerning the first ultraintelligent machine. *Advances in computers* (Vol. 6, pp. 31–88). Elsevier.
Grace, K., Salvatier, J., Dafoe, A., Zhang, B., & Evans, O. (2018). When will AI exceed human performance? Evidence from AI experts. *Journal of Artificial Intelligence Research*, *62*, 729–754.
Grace, K., Stewart, H., Sandkühler, J. F., Thomas, S., Weinstein-Raun, B., & Brauner, J. (2024). Thousands of AI authors on the future of AI. *arXiv preprint arXiv:2401.02843*.
Greenblatt, R., Shlegeris, B., Sachan, K., & Roger, F. (2023). Ai control: Improving safety despite intentional subversion. *arXiv preprint arXiv:2312.06942*.
Häggström, O. (2016). *Here be dragons: Science, technology and the future of humanity*. Oxford University Press.
Häggström, O. (2018). Challenges to the Omohundro–Bostrom framework for AI motivations. *Foresight*, *21*(1), 153–166.
Huang, C., Zhang, Z., Mao, B., & Yao, X. (2022). An overview of artificial intelligence ethics. *IEEE Transactions on Artificial Intelligence*, *4*(4), 799–819.
Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., Gray, S., Radford, A., Wu, J., & Amodei, D. (2020). Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*.
Katz, D. M., Bommarito, M. J., Gao, S., & Arredondo, P. (2024). Gpt-4 passes the bar exam. *Philosophical Transactions of the Royal Society A*, *382*(2270), 20230254.
Kohlberg, L., & Hersh, R. H. (1977). Moral development: A review of the theory. *Theory into Practice*, *16*(2), 53–59.
Kurzweil, R. (1990). *The age of intelligent machines* (Vol. 580). MIT Press.
Kurzweil, R. (2000). *The age of spiritual machines: When computers exceed human intelligence*. Penguin.
Kurzweil, R. (2005). The singularity is near. *Ethics and emerging technologies* (pp. 393–406). Springer.
Landgrebe, J., & Smith, B. (2022). *Why machines will never rule the world: Artificial intelligence without fear*. Routledge.
Larsson, S., & Heintz, F. (2020). Transparency in artificial intelligence. *Internet Policy Review*, *9*(2).
Legg, S., & Hutter, M. (2007). A collection of definitions of intelligence. *Frontiers in Artificial Intelligence and Applications*, *157*, 17.
Lloyd, S. (2000). Ultimate physical limits to computation. *Nature*, *406*(6799), 1047–1054.
Lucas, J. R. (1961). Minds, machines and Gödel. *Etica E Politica*, *5*(1), 1.

Mankowitz, D. J., Michi, A., Zhernov, A., Gelmi, M., Selvi, M., Paduraru, C., Leurent, E., Iqbal, S., Lespiau, J. B., & Ahern, A. (2023). Faster sorting algorithms discovered using deep reinforcement learning. *Nature*, *618*(7964), 257–263.

Meyer, A., Riese, J., & Streichert, T. (2024). Comparison of the performance of GPT-3.5 and GPT-4 with that of medical students on the written German medical licensing examination: Observational study. *JMIR Medical Education*, *10*, e50965.

Miller, J. D., Yampolskiy, R., & Häggström, O. (2020). An AGI modifying its utility function in violation of the strong orthogonality thesis. *Philosophies*, *5*(4), 40.

Mitchell, S., Potash, E., Barocas, S., D'Amour, A., & Lum, K. (2021). Algorithmic fairness: Choices, assumptions, and definitions. *Annual Review of Statistics and its Application*, *8*(1), 141–163. https://doi.org/10.1146/annurev-statistics-042720-125902

Morris, M. R., Sohl-Dickstein, J., Fiedel, N., Warkentin, T., Dafoe, A., Faust, A., Farabet, C., & Legg, S. (2023). Levels of AGI: Operationalizing Progress on the Path to AGI. *arXiv preprint arXiv:2311.02462*.

Müller, V. C., & Bostrom, N. (2016). Future progress in artificial intelligence: A survey of expert opinion. *Fundamental Issues of Artificial Intelligence*, 555–572.

Müller, V. C., & Cannon, M. (2022). Existential risk from AI and orthogonality: Can we have it both ways? *Ratio*, *35*(1), 25–36.

Nowak, M. A. (2006). Five rules for the evolution of Cooperation. *Science*, *314*(5805), 1560–1563.

Penrose, R. (1994). *Shadows of the mind: A search for the missing science of consciousness*. Oxford University Press.

Railton, P. (2020). Ethical learning, natural and artificial. *Ethics of Artificial Intelligence*, *45*, 78.

Russell, S. (2019). *Human compatible: Artificial intelligence and the problem of control*. Penguin.

Searle, J. R. (1980). Minds, brains, and programs. *Behavioral and Brain Sciences*, *3*(3), 417–424.

Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askell, A., Bowman, S. R., Cheng, N., Durmus, E., Hatfield-Dodds, Z., & Johnston, S. R. (2023). Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*.

Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., & Graepel, T. (2018). A general reinforcement learning algorithm that masters chess, shogi, and go through self-play. *Science*, *362*(6419), 1140–1144.

Soori, M., Arezoo, B., & Dastres, R. (2023). Artificial intelligence, machine learning and deep learning in advanced robotics, a review. *Cognitive Robotics*, *3*, 54–70. https://doi.org/10.1016/j.cogr.2023.04.001

Sparrow, R. (2024). Friendly AI will still be our master. Or, why we should not want to be the pets of super-intelligent computers. *AI & SOCIETY*, *39*(5), 2439–2444. https://doi.org/10.1007/s00146-023-01698-x

Stenseke, J. (2024). On the computational complexity of ethics: Moral tractability for Minds and machines. *Artificial Intelligence Review*, *57*(4), 105. https://doi.org/10.1007/s10462-024-10732-3

Tegmark, M. (2017). *Life 3.0: Being human in the age of artificial intelligence*. Knopf.

Thorstad, D. (2024). Against the singularity hypothesis. *Philosophical Studies*. https://doi.org/10.1007/s11098-024-02143-5

Tubert, A., & Tiehen, J. (2024). Existentialist risk and value misalignment. *Philosophical Studies*. https://doi.org/10.1007/s11098-024-02142-6

Vallor, S. (2024). *The AI mirror: How to reclaim our humanity in an age of machine thinking*. Oxford University Press.

Van Wynsberghe, A. (2021). Sustainable AI: AI for sustainability and the sustainability of AI. *AI and Ethics*, *1*(3), 213–218.

Vinge, V. (1993). The coming technological singularity: How to survive in the post-human era. *Science Fiction Criticism: an Anthology of Essential Writings*, *81*, 352–363.

Yampolskiy, R. V. (2012). Leakproofing the singularity artificial intelligence confinement problem. *Journal of Consciousness Studies JCS*.

Yampolskiy, R. V. (2013). Artificial intelligence safety engineering: Why machine ethics is a wrong approach. *Philosophy and theory of artificial intelligence* (pp. 389–396). Springer.

Yampolskiy, R. V. (2024). *AI: Unexplainable, unpredictable, uncontrollable*. CRC.

Yudkowsky, E. (2008). Artificial intelligence as a positive and negative factor in global risk. *Global Catastrophic Risks*, *1*(303), 184.

**Publisher's note** Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.